

Serving Tensorflow Models with Kubernetes

Ron Bodkin, Technical Director, Applied AI

Brian Foo, Senior Software Engineer, Applied AI

Step One: Get An Account
Go To goo.gl/sqZ2Qp to start

Google Cloud



Your Instructors



Ron
Bodkin



Brian
Foo




Cassie
Kozyrkov

Set up an account here: goo.gl/sgZ2Qp

Your Instructors



Ron
Bodkin



Brian
Foo



Cassie
Kozyrkov

Set up an account here: goo.gl/sgZ2Qp

Overview

- Tensorflow
 - Building and Running Tensorflow graphs & APIs
 - Tensorflow Serving for Online Predictions
- Docker and Kubernetes
 - A brief history of containers
 - Kubernetes: pods and services
- The Marriage of TF Serving and Kubernetes



Building and Running Tensorflow Graphs



Build a Tensorflow Graph

```
import tensorflow as tf

x = tf.placeholder(tf.float32, shape=(100))

# Some preloaded model weights and biases
w = tf.get_variable('weights', shape=(100))
b = tf.get_variable('bias', shape=[])
y = tf.tensordot(w, x, 1) + b
```



Run the Tensorflow Graph

```
import tensorflow as tf
```

```
x = tf.placeholder(tf.float32, shape=(100))
```

```
# Some preloaded model weights and biases
```

```
w = tf.get_variable('weights', shape=(100))
```


```
b = tf.get_variable('bias', shape=[])
```

```
y = tf.tensordot(w, x, 1) + b
```

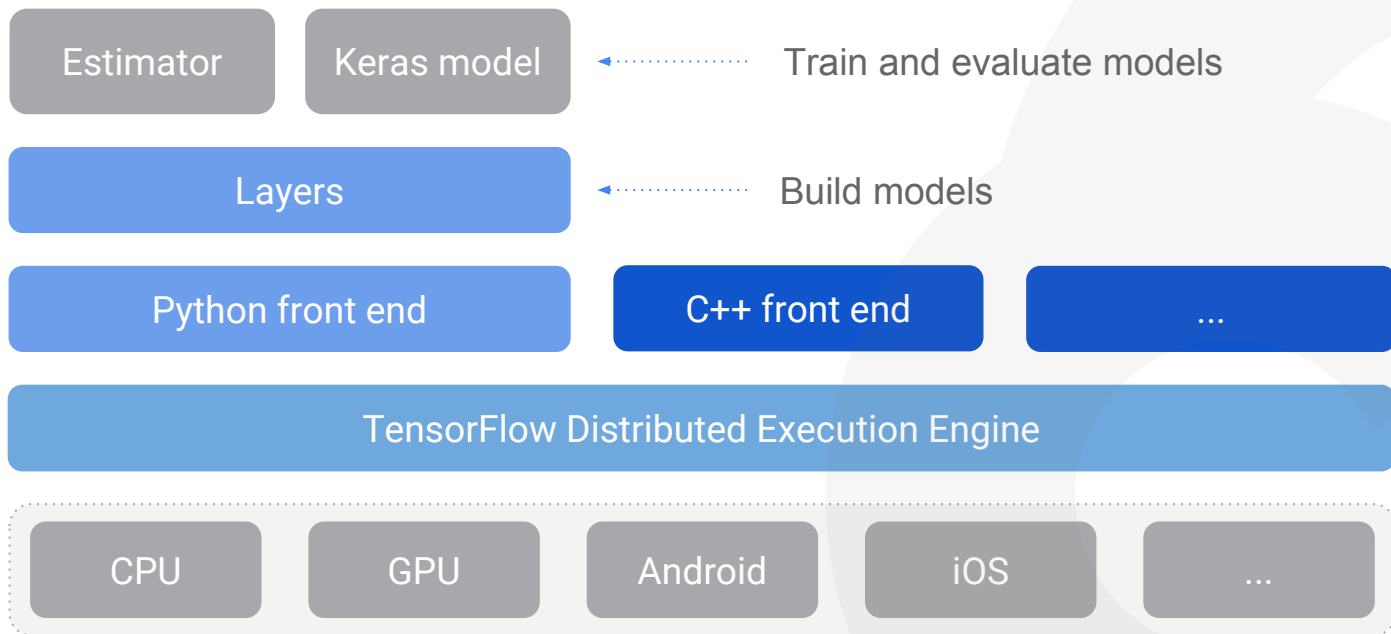
```
with tf.Session() as sess:
```

```
    print(sess.run(y, feed_dict={x: input_data}))
```

Data entrypoint for predictions



Different APIs for TF




What is Serving?


Serving is how you *apply* a ML model,
after you've trained it



What is Serving?



What is Serving?

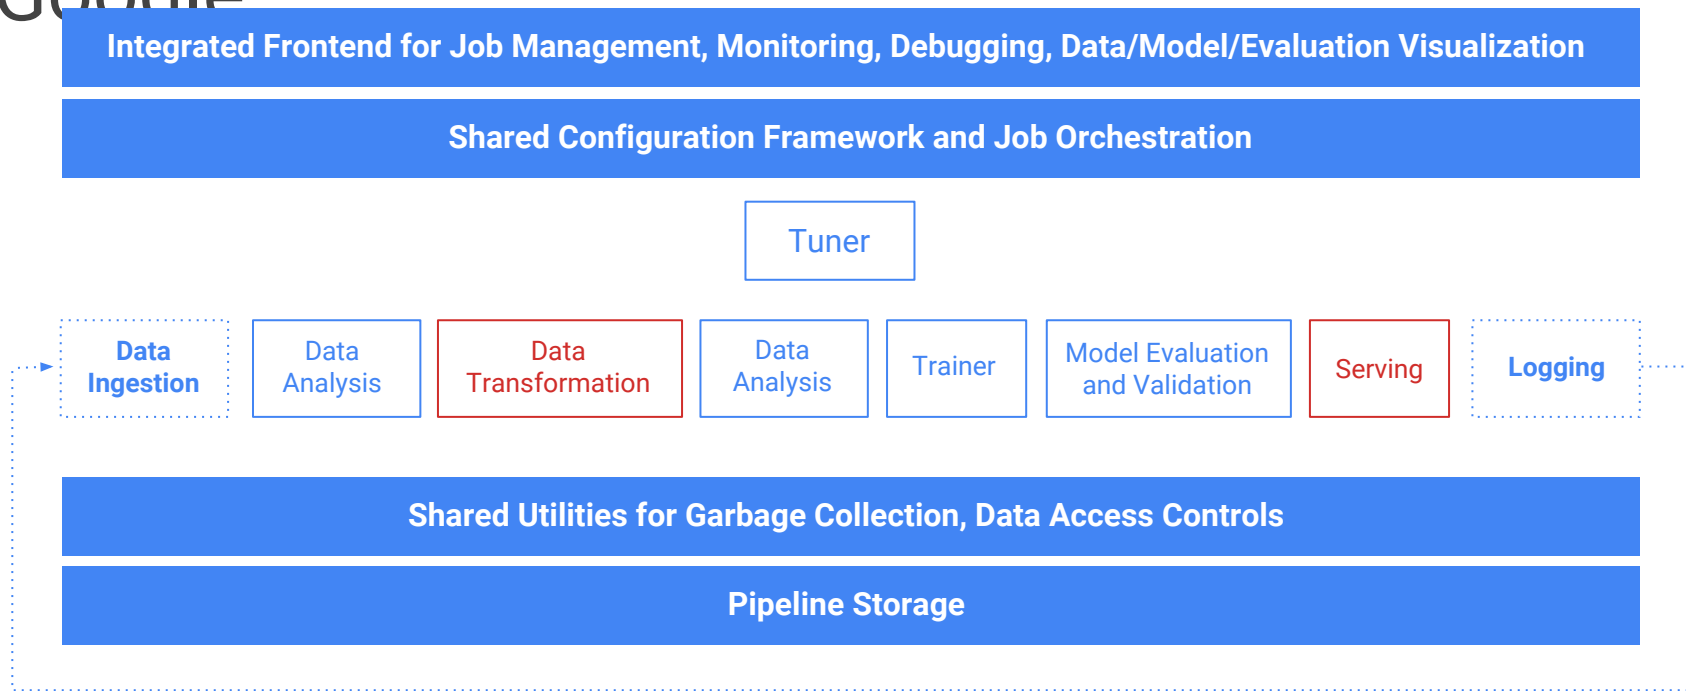


TensorFlow Serving Goals

- Online, low latency
- Multiple models in a single process
- Multiple *versions* of a model loaded over time
- Compute cost varies in real-time to meet product demand
 - auto-scale with CloudML, Docker & K8s
- Aim for the efficiency of mini-batching at training time ...
 - except with requests arriving asynchronously




TFX: TensorFlow-Based ML deployment at Google



Docker and Kubernetes



A Brief History of Containers




2008

In 2008, Linux introduced containers.

- Isolated environment for running applications, except...
- All applications must have a common OS Kernel (e.g. Ubuntu, Debian, etc.)

Enter Docker...



2008

2013

In 2013, Docker found a way around the shared kernel problem.

- Application containers have their own OS kernel.
- Flexible resource requirements.
- Perfect for cloud computing!

Enter Kubernetes



2008


2013

2014

In 2014, Google open sourced Kubernetes.

- Deploy Docker containers to any number of machines
- Create load balancing and front-end services to handle external requests.
- Automatically restart backend containers when they fail.

And Many More...



Heapster, KSonnet, Kubeflow, ...

2008

2013

2014

2015...


Example projects built on top of Kubernetes:

- Monitoring (Heapster)
- Deployment languages (KSonnet)
- Deployment automation (Kubeflow for ML)
- and many more!

Kubernetes in a Nutshell

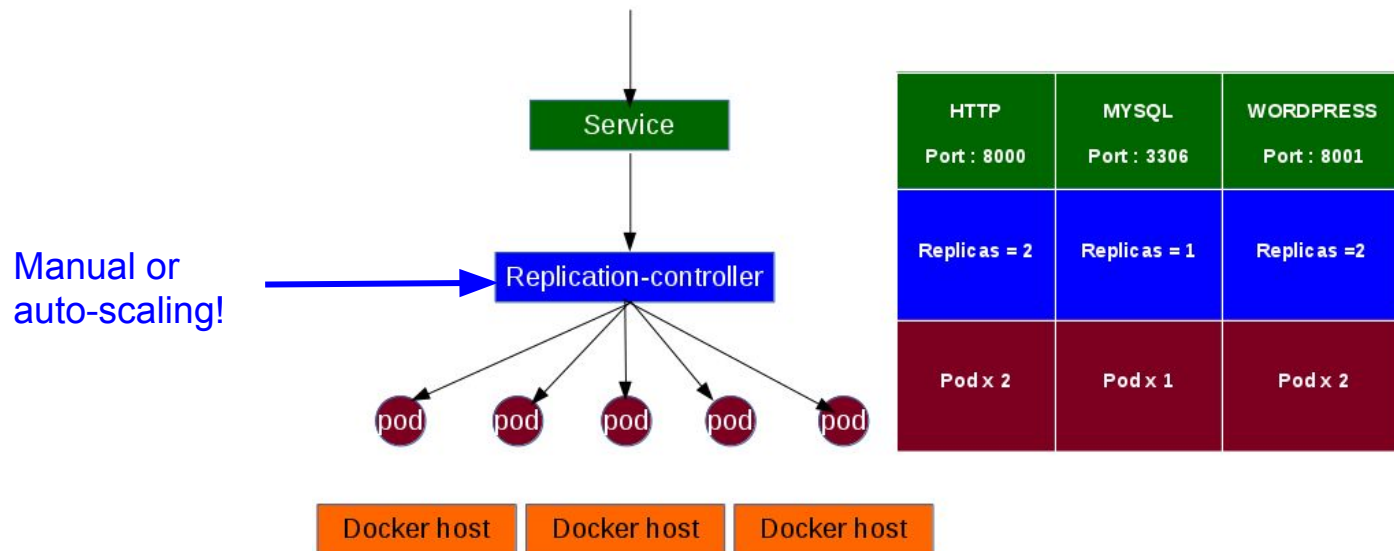
Kubernetes Pods

collections of containers that are co-scheduled



Kubernetes in a Nutshell

How Kubernetes Works?



Tensorflow Serving using Kubernetes



TF Serving on Kubernetes Workflow




→ Label: cat
Prob: 97.6%

What do we want?

- A prediction service that can handle multiple client requests
- Load-balancing across TF model servers
- Ability to scale up

TF Serving on Kubernetes Workflow

Exercises



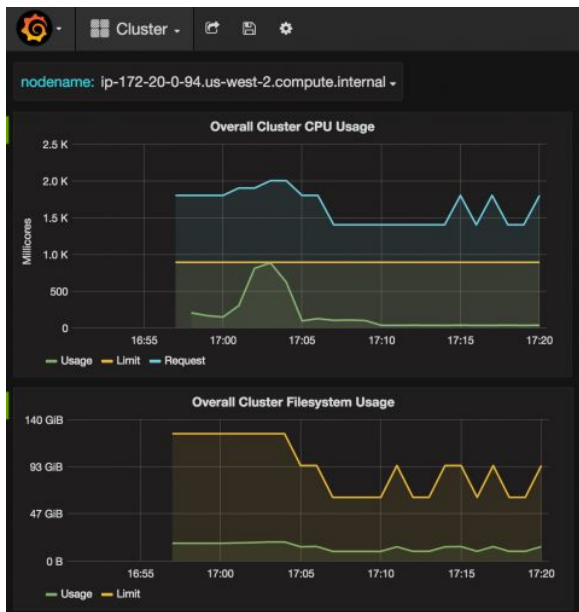
How do we get there?

- Convert TF training code to model for serving.
- Package model in Docker container and upload to a registry.
- Use **Kubernetes** to:
 - Deploy container on multiple back-end pods.
 - Deploy a front-end service to send client requests to a backend pod.
- Send protobufs (encoded JSONs) containing images to kubernetes cluster.
- Load testing.

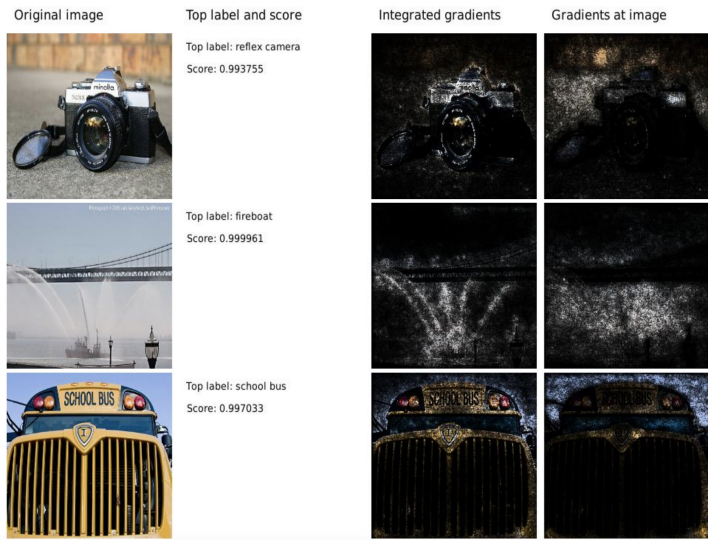
Monitoring, Interpretation, and Keras

Bonus Exercises

Heapster Grafana Dashboard (Pod and Cluster Resource Usage)



Model Understanding and Visualization using Integrated Gradients



Codelab Time!

Open a Chrome incognito window.

Log in at events.qwiklabs.com

If you don't have an account register at goo.gl/sgZ2Qp



Recap and Demos



Acknowledgements and Additional Resources

Special thanks to:

- [Kubeflow](#): providing Docker images and templates for TF Serving on Kubernetes
David Aronchick, Jeremy Lewi, Vishnu Kannan (Google); Peng Yu (Shopify)
- [Google Cloud ML](#): GPU batch profiling work using Beam and Tensorflow

Reference - Model Visualization:

- Sundararajan, Taly, Yan. ***Axiomatic Attribution for Deep Networks***. ICML 2017. Link: arxiv.org/pdf/1703.01365.pdf

Thank you

1. Please leave feedback
2. Resources at goo.gl/Sg6ecA
3. Save any work you want to keep

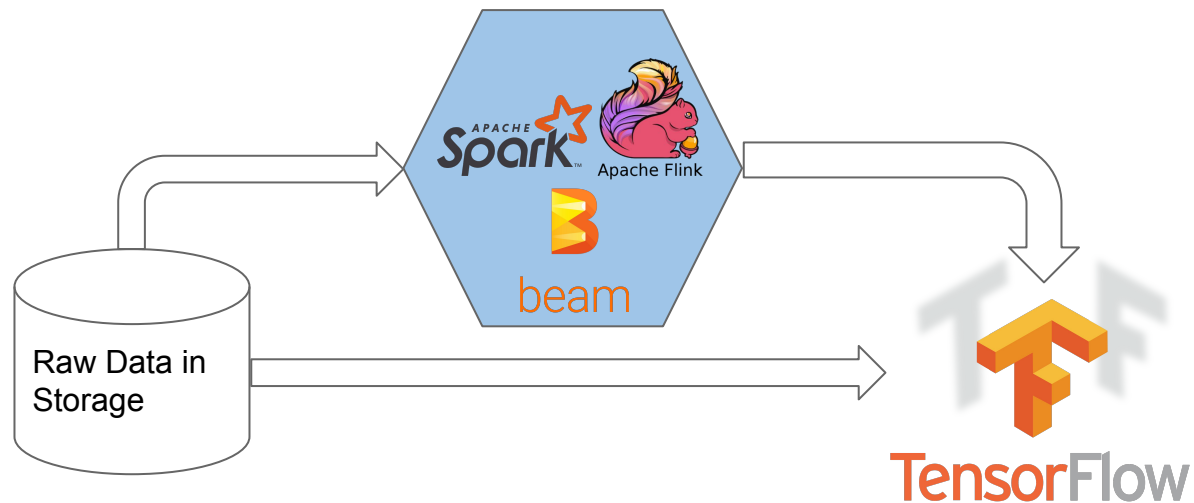


Appendix: pipelines versus Client-Server Architectures



Pipeline Architecture: Batch/Online Processing

- Read offline data from local/HDFS/Google Storage/AWS
- Preprocess (clean, filter, aggregate) using Spark/Beam/Flink
- Create batches to run through a TF graph
- Update model params (training) / Collect inference results (serving)



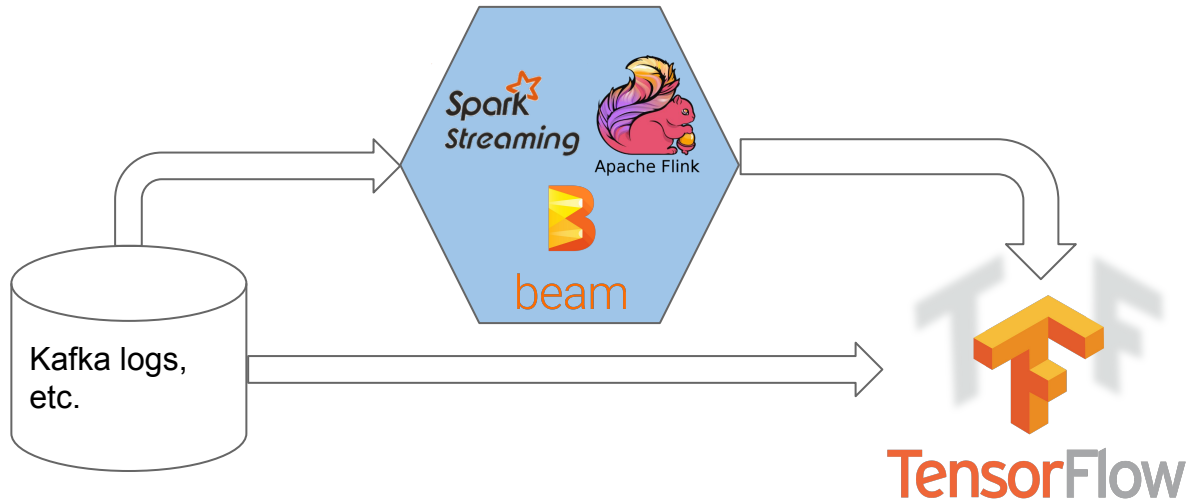
Pipeline Architecture: Batch/Online Processing

- **Benefits:**

- Full control over pipeline application and model!


- **Limitations:**

- **Language Dependency:** Requires Python, or Java JNI to C++
- **No Proprietary Models:** Requires graph and model params to be exposed in code.
- **Experience:** Months to years of expertise to build, debug, and manage pipelines effectively.




Client-Server Architecture: Tensorflow Serving

- Asynchronous and Streaming Model Serving
- Efficient implementation in c++
- Server build can be optimized for native environment
 - CPUs or GPUs
 - Just-in-time (JIT) compilation
 - etc.



Client-Server Architecture: Tensorflow Serving

- Asynchronous and Streaming Model Serving
- Efficient implementation in c++
- Build can be optimized for the environment (CPUs or GPUs)
- **Language independent Protobufs!**
 - RESTful API calls using serialized dictionaries
 - Send dictionary of data
 - Receive dictionary of prediction results



Client-Server Architecture: Tensorflow Serving

- Asynchronous and Streaming Model Serving
- Efficient implementation in c++
- Build can be optimized for the environment (CPUs or GPUs)
- **Language independent Protobufs!**
 - RESTful API calls using serialized dictionaries
 - Send dictionary of data
 - Receive dictionary of prediction results
- How do we guarantee identical serving environments?
- How do we scale?
- How do we handle failures gracefully?

