



How to build leakproof stream processing pipelines with Apache Kafka and Apache Spark

Introduction

- Guru Medasani
 - Data Science Architect at Domino Data Lab
 - Previously senior solutions architect at Cloudera
- Jordan Hambleton - Consulting Manager in San Francisco
 - Nearly 4 years as Resident Senior Architect at large technology firm
 - Previously software engineer building operational data systems on CDH

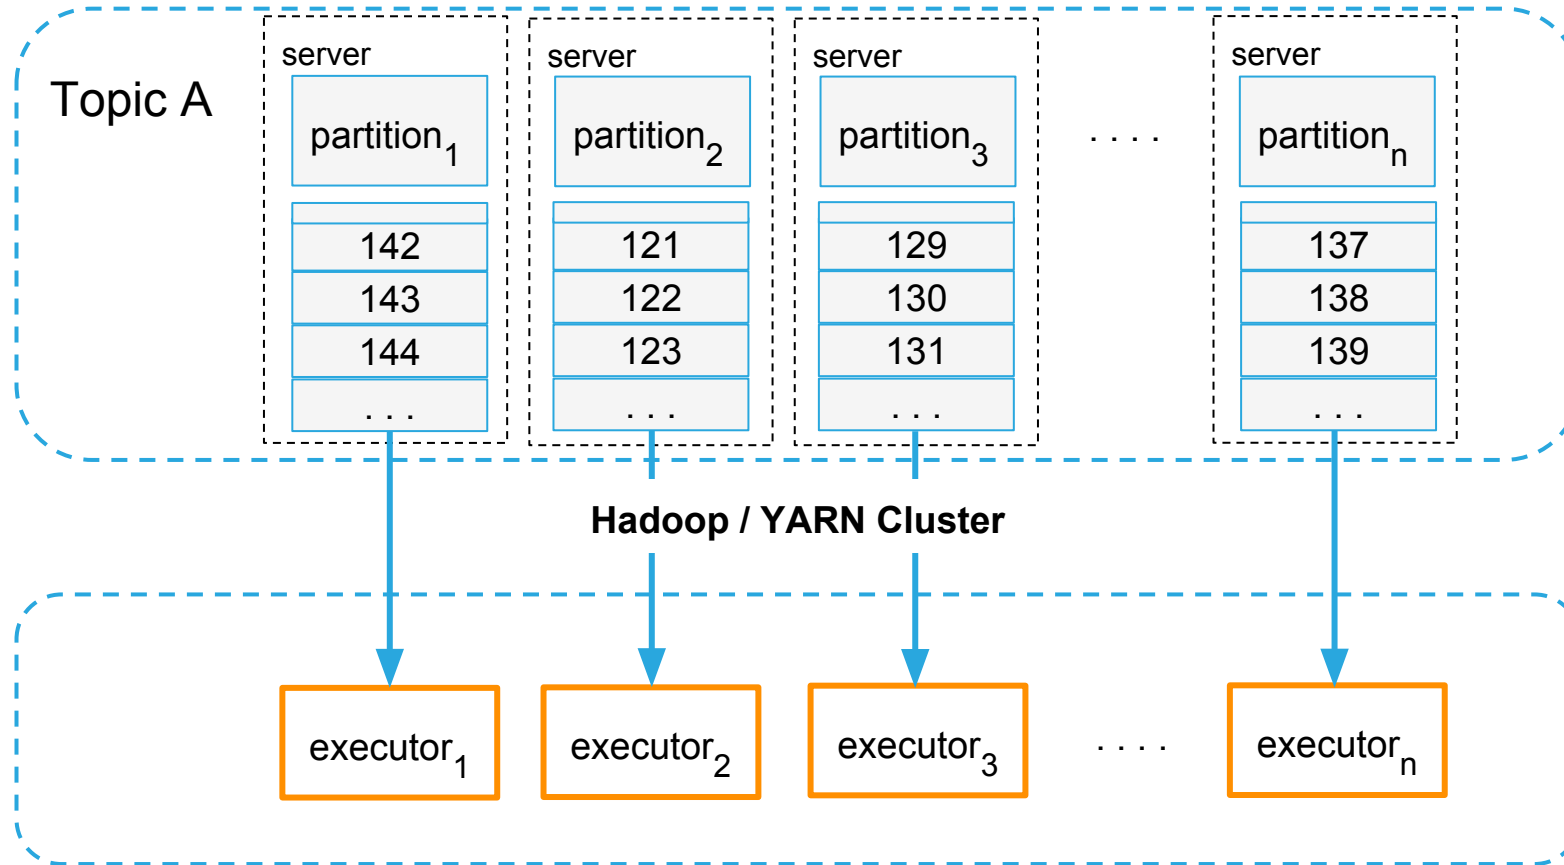
Agenda

- Intro
- Overview of Spark Streaming from Kafka
 - Workflow of the DStream and RDD
 - Spark Streaming Kafka consumer types
- Offset management
 - Motivation
 - Storing offsets in external data stores
- Q & A




Overview

Kafka Cluster



more parallelism

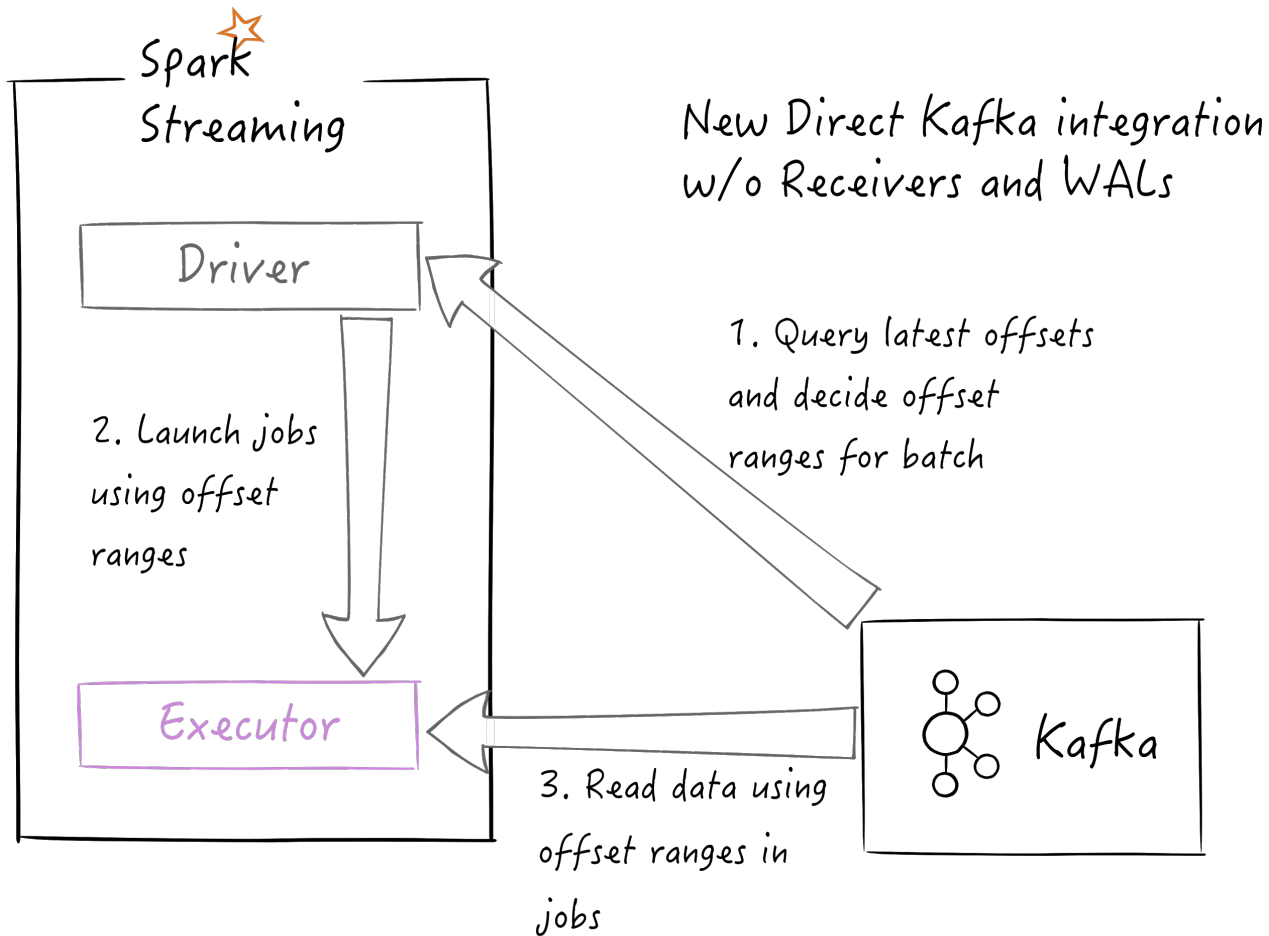


Spark
Streaming

Overview Spark Streaming from Kafka

- DStream - sequence of RDDs
- Two approaches in KafkaUtils
 - Receiver based
 - Direct approach (recommended & the method we talk about)
- Spark streaming embeds a kafka client
 - Spark 1.6 uses the 0.9.0-kafka-2.0.0 client (SimpleConsumer)
 - Spark 2.x kafka 0-8-0 uses the 0.9.0-kafka-2.0.2 client (SimpleConsumer)
 - Spark 2.x kafka 0-10-0 uses the 0.10.0-kafka-2.1.0 client (KafkaConsumer)

DStream and RDD Workflow



- Spark Streaming
 - `batchIntervalInSeconds`
 - `stopGracefullyOnShutdown`
- Kafka
 - `bootstrap.servers`
 - `auto.offset.reset`
 - `group.id`
 - `key.deserializer`
 - `value.deserializer`

Spark Streaming Kafka Consumer # 1

- spark-streaming-kafka-0-8 / 0.9.0-kafka-2.0.2
- DStream
 - Gets range of each topic/partition - throttle maxRatePerPartition
 - auto.offset.reset (smallest|largest)
 - refresh.leader.backoff.ms - lost leader
- KafkaRDD for set of topic, partition, offsets
 - User can now get offset ranges from RDD
 - topic, partition, fromOffset (inclusive), untilOffset (exclusive)
- KafkaRDDPartition iterator
 - SimpleConsumer initialized and batches of events fetched
 - refresh.leader.backoff.ms - lost leader


Spark Streaming Kafka Consumer # 2

- **Supported** - spark-streaming-kafka-0-10 / 0.10.0-kafka-2.1.0
- Internal Kafka client uses new Java KafkaConsumer
- ConsumerStrategies
 - subscribe, assign, subscribe pattern
- LocationStrategies
 - executor distribution strategy (consistent, fixed, brokers)
- DStream
 - Gets range of each topic/partition - throttle maxRatePerPartition
 - auto.offset.reset (earliest|latest)
 - Be careful - enable.auto.commit (default true)
 - heartbeat & session timeouts

Spark Streaming Kafka Consumer # 2

- DStream
 - Consumer poll for group coordination & discovery
 - Identify new partitions, from offsets
 - Pause consumer
 - seekToEnd to get untilOffsets
- KafkaRDD
 - Fixed [enable.auto.commit = false, auto.offset.reset = none, spark-executor-`{group.id}`]
 - Attempts to assign offset range consistently for optimal consumer caching
- KafkaRDDPartition iterator
 - Initialize/lookup CachedKafkaConsumer with executor group
 - consumer assigned per single topic, partition with internal buffer
 - on cache miss, seek and poll

Keeping Track



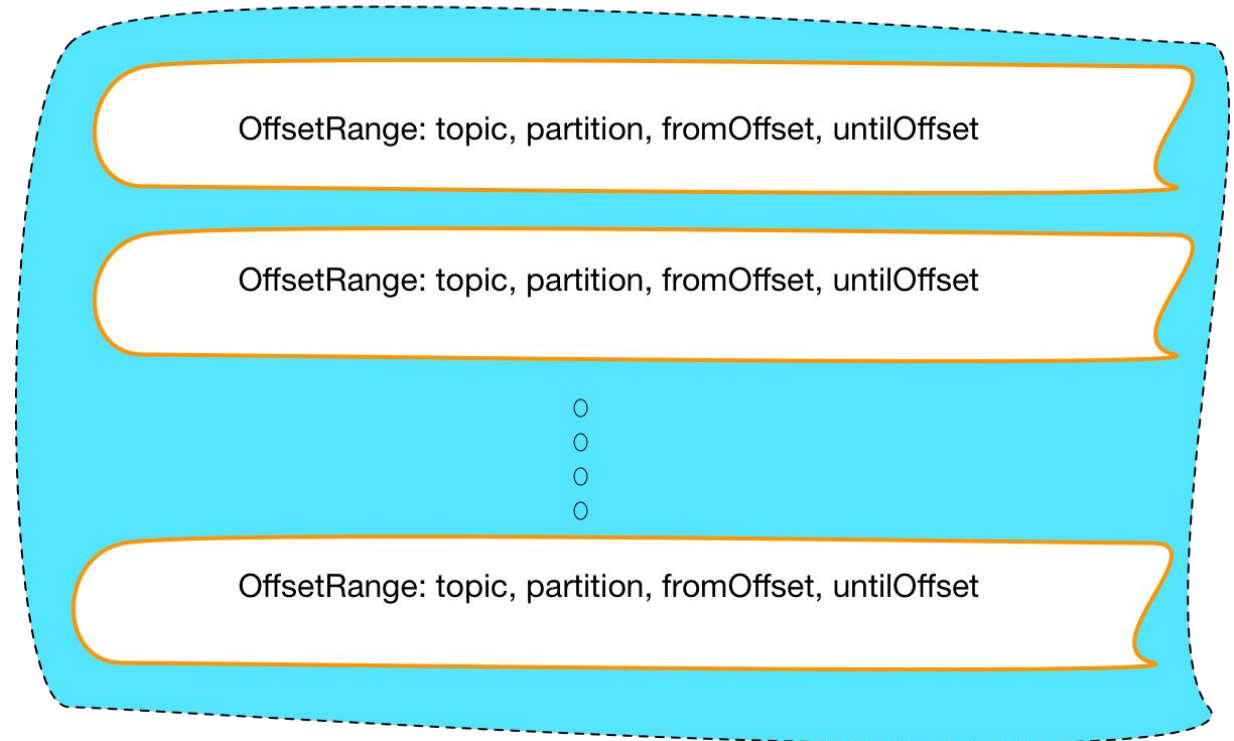
Motivation for Tracking Offsets

- Planned Maintenance
 - Upgrades
 - Bug-fixes
- Unplanned Maintenance
 - Failures
- Application Processing Errors
 - Wrong calculations
 - Updated algorithm over known streaming data
- More control over messages
 - Just earliest and latest are insufficient

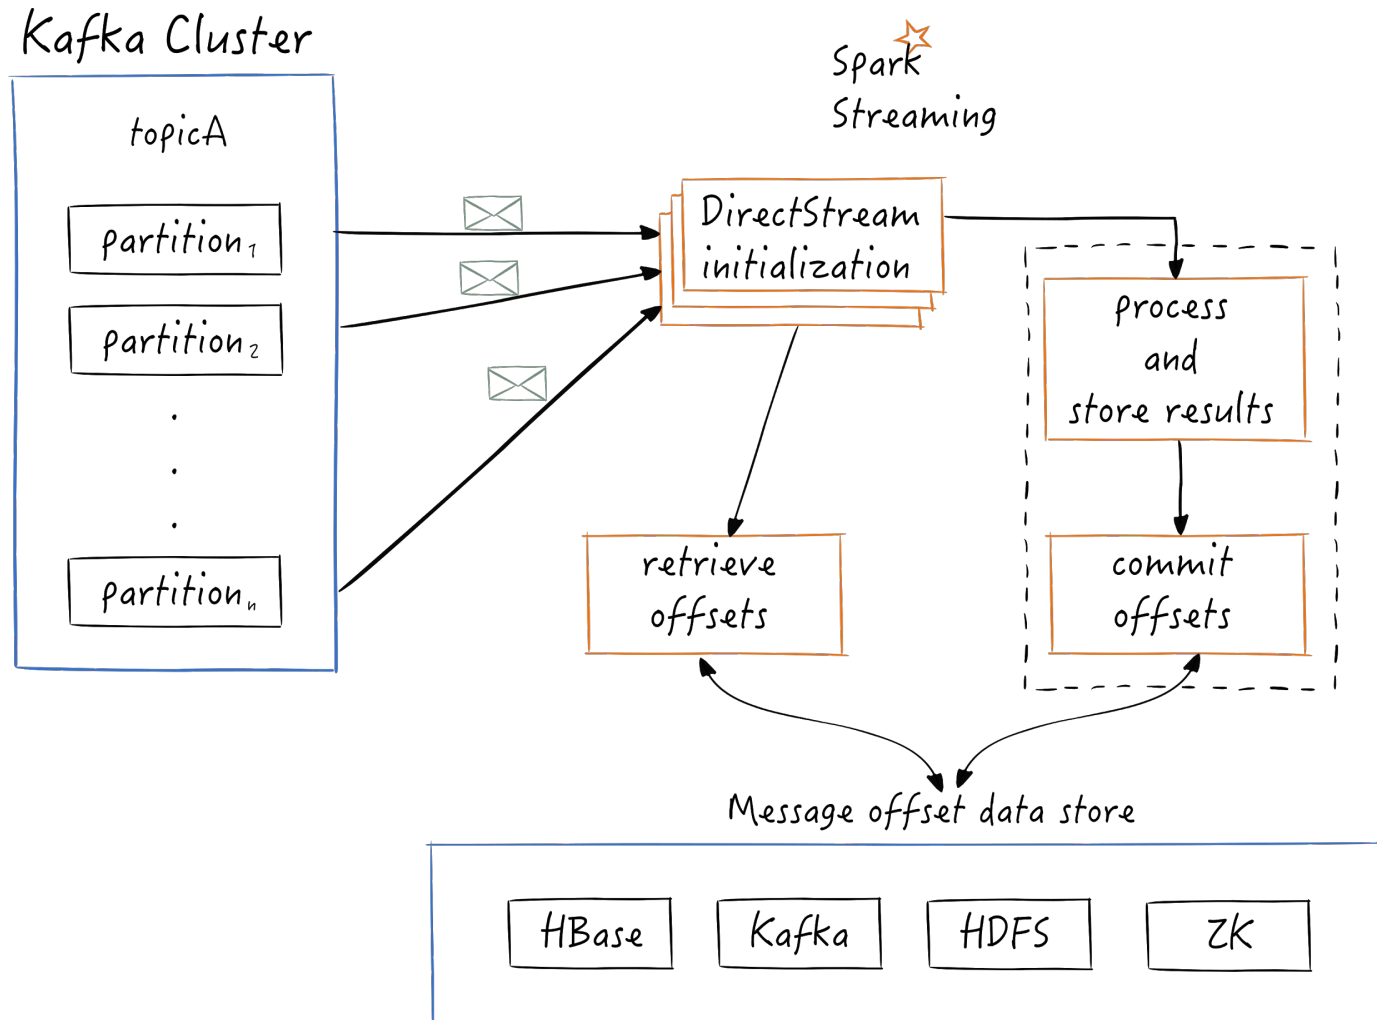
Obtaining Offsets

- Cast RDD to HasOffsetRanges
- DStream's first transformation

Array[OffsetRange]



Offset management Workflow



- Limited options prior to spark-streaming-kafka-0-10
- Store offsets in external datastore
 - Checkpoints (Not recommended)
 - ZooKeeper
 - Kafka
 - HBase
- Do not have to manage offsets

Offset Management in ZooKeeper

- ZooKeeper
 - znode - /consumers/[groupId]/offsets/[topic]/[partitionId] -> long (offset)
 - Only retains latest committed offsets
 - Can easily be managed by external tools
 - Leverage existing monitoring for Lag, no historical insight

Offset Management in Kafka

- Kafka
 - CanCommitOffsets provides async commit to internal kafka topic
 - More difficult to manage internal kafka topic manually
 - Leverage existing monitoring for Lag, no historical insight

Offset Management in HBase

- HBase
 - Unique entry per consumer group, batch

schema:

row:	<TOPIC_NAME>:<GROUP_ID>:<EPOCH_BATCHTIME_MS>
column family:	offsets
qualifier:	<PARTITION_ID>
value:	<OFFSET_ID>

- Fine-grained monitoring over time
- HBase shell for easy management
- Get latest entry -
 - `scan 'prod_stream',`
 - `STARTROW =>'device_alerts:csi_group',`
 - `REVERSED =>TRUE,`
 - `LIMIT =>1`

```
hbase(main):001:0> scan 'stream_kafka_offsets', {REVERSED => true}
ROW
kafkablog2:groupid-1:1497628830000    column=offsets:0, timestamp=1497628832448, value=285
kafkablog2:groupid-1:1497628830000    column=offsets:1, timestamp=1497628832448, value=285
kafkablog2:groupid-1:1497628830000    column=offsets:2, timestamp=1497628832448, value=285
kafkablog2:groupid-1:1497628770000    column=offsets:0, timestamp=1497628773773, value=225
kafkablog2:groupid-1:1497628770000    column=offsets:1, timestamp=1497628773773, value=225
kafkablog2:groupid-1:1497628770000    column=offsets:2, timestamp=1497628773773, value=225
kafkablog1:groupid-2:1497628650000    column=offsets:0, timestamp=1497628653451, value=165
kafkablog1:groupid-2:1497628650000    column=offsets:1, timestamp=1497628653451, value=165
kafkablog1:groupid-2:1497628650000    column=offsets:2, timestamp=1497628653451, value=165
kafkablog1:groupid-1:1497628530000    column=offsets:0, timestamp=1497628533108, value=120
kafkablog1:groupid-1:1497628530000    column=offsets:1, timestamp=1497628533108, value=120
kafkablog1:groupid-1:1497628530000    column=offsets:2, timestamp=1497628533108, value=120
4 row(s) in 0.5030 seconds

hbase(main):002:0>
```


Starting Streaming Jobs with Known Offsets

- Spark Streaming job started for the first time
- No changes in Kafka partitions
- Increase in number of Kafka partitions

<http://blog.cloudera.com/blog/2017/06/offset-management-for-apache-kafka-with-apache-spark-streaming/>

Questions?

Thank you

Jordan Hambleton
Guru Medasani

